

# Open Neighbourhoods Precede Concept Spread: Structural Antecedents and Cross-Disciplinary Diffusion of Emerging Scientific Concepts

Internal Research Report

September 2026

## Contents

|          |                                                                                |          |
|----------|--------------------------------------------------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                                                            | <b>4</b> |
| <b>2</b> | <b>Iteration 1: Data Foundation and Pre-Registration</b>                       | <b>4</b> |
| 2.1      | Strategy . . . . .                                                             | 4        |
| 2.2      | Artifact 1: Concept Pool and OpenAlex Work Corpus . . . . .                    | 5        |
| 2.3      | Gate A: Citation-Layer Feasibility . . . . .                                   | 5        |
| 2.4      | Denominators and Venue Habitat . . . . .                                       | 6        |
| 2.5      | Artifact 2: Whole-Science Background Sample . . . . .                          | 6        |
| 2.6      | Artifact 3: Held-Out MeSH Confirmation Population . . . . .                    | 6        |
| 2.7      | Artifact 4: Semantic Grounding and Labelling Infrastructure . . . . .          | 7        |
| 2.8      | Artifact 5: Prior-Art Dossier and Pre-Registration . . . . .                   | 7        |
| 2.9      | Dead Ends and Shortfalls (Iteration 1) . . . . .                               | 7        |
| 2.10     | What Was Learned . . . . .                                                     | 8        |
| <b>3</b> | <b>Iteration 2: Network Construction, Gate A Failure, and the Closure Lead</b> | <b>8</b> |
| 3.1      | Strategy . . . . .                                                             | 8        |
| 3.2      | Artifact 6: Hydrated Corpus, Nativeness Profiles, and Venue Habitat . . . . .  | 9        |
| 3.3      | Artifact 7: Concept Grounding Pipeline . . . . .                               | 9        |
| 3.4      | Artifact 8: Viability Layer, Gate A, Viability States . . . . .                | 9        |
| 3.4.1    | Gate A Re-Test at Edge Level . . . . .                                         | 10       |
| 3.4.2    | Viability States (Descriptive Only) . . . . .                                  | 11       |
| 3.4.3    | Power Before Any Outcome . . . . .                                             | 11       |
| 3.5      | Artifact 9: Precursor Event Study, Main Pool . . . . .                         | 11       |
| 3.5.1    | Emergence Labels . . . . .                                                     | 11       |
| 3.5.2    | Event Study Results . . . . .                                                  | 11       |
| 3.5.3    | Prediction . . . . .                                                           | 12       |
| 3.5.4    | Trajectory Patterns and Typology . . . . .                                     | 12       |
| 3.6      | Artifact 10: Precursor Replication, MeSH Population . . . . .                  | 12       |
| 3.7      | Decision-Rule Evaluation . . . . .                                             | 12       |
| 3.8      | Dead Ends and Negative Results (Iteration 2) . . . . .                         | 13       |
| 3.9      | What Was Learned . . . . .                                                     | 13       |

|          |                                                                                |           |
|----------|--------------------------------------------------------------------------------|-----------|
| <b>4</b> | <b>Iteration 3: Deepening the Closure Lead and the Grafting Test</b>           | <b>13</b> |
| 4.1      | Strategy . . . . .                                                             | 13        |
| 4.2      | Artifact 11: Record Audit and Turnover Pre-Check . . . . .                     | 13        |
| 4.3      | Artifact 12: Openness Before Take-Off . . . . .                                | 14        |
| 4.4      | Artifact 13: Breadth-Prediction Screen . . . . .                               | 15        |
| 4.5      | Artifact 14: Host-Entry Grafting Test . . . . .                                | 15        |
| 4.6      | Artifact 15: Descriptive Diffusion . . . . .                                   | 16        |
|          | 4.6.1 Diffusion Typology . . . . .                                             | 16        |
|          | 4.6.2 Community Roles . . . . .                                                | 16        |
|          | 4.6.3 Expansion Precedes Diffusion . . . . .                                   | 16        |
| 4.7      | Dead Ends and Negative Results (Iteration 3) . . . . .                         | 16        |
| 4.8      | What Was Learned . . . . .                                                     | 18        |
| <b>5</b> | <b>Iteration 4: Confirmation on Held-Out Folds</b>                             | <b>18</b> |
| 5.1      | Strategy . . . . .                                                             | 18        |
| 5.2      | Reviewer Critique Disposition . . . . .                                        | 18        |
| 5.3      | Artifact 16: Held-Out Grafting Confirmation and Discriminating Tests . . . . . | 18        |
|          | 5.3.1 Held-Out Grafting Confirmation . . . . .                                 | 18        |
|          | 5.3.2 Single-Paper Artefact Test . . . . .                                     | 20        |
|          | 5.3.3 Host-Vocabulary Decomposition . . . . .                                  | 20        |
|          | 5.3.4 Topic-Score Absorption . . . . .                                         | 20        |
| 5.4      | Artifact 17: MeSH Replication of Grafting . . . . .                            | 20        |
| 5.5      | Artifact 18: Closure Held-Out Confirmation . . . . .                           | 22        |
|          | 5.5.1 Closure-Anchoring Link Test . . . . .                                    | 22        |
| 5.6      | Artifact 19: Descriptive Held-Out Confirmation . . . . .                       | 22        |
|          | 5.6.1 Lead-Lag: Expansion Before Diffusion . . . . .                           | 22        |
|          | 5.6.2 Held-Out Confirmation of Descriptive Findings . . . . .                  | 22        |
|          | 5.6.3 Rooting . . . . .                                                        | 22        |
| 5.7      | Artifact 20: Adopter-Level Mechanism . . . . .                                 | 23        |
| 5.8      | Dead Ends and Negative Results (Iteration 4) . . . . .                         | 23        |
| 5.9      | What Was Learned . . . . .                                                     | 24        |
| <b>6</b> | <b>Iteration 5: Post-Confirmation Decomposition and Full Audit</b>             | <b>24</b> |
| 6.1      | Strategy . . . . .                                                             | 24        |
| 6.2      | Artifact 21: Margin Decomposition and Field-Boundary Test . . . . .            | 24        |
|          | 6.2.1 Extensive vs Intensive Margin . . . . .                                  | 24        |
|          | 6.2.2 Field Boundary . . . . .                                                 | 26        |
| 6.3      | Artifact 22: Host-Specificity Test . . . . .                                   | 26        |
| 6.4      | Artifact 23: Full Numerical Audit . . . . .                                    | 26        |
|          | 6.4.1 Table 18 Self-Test Honesty . . . . .                                     | 27        |
| 6.5      | Representative-Case Interpretations . . . . .                                  | 27        |
| 6.6      | Artifact 24: Positioning Dossier . . . . .                                     | 27        |
| 6.7      | Artifact 25: Paper Figures . . . . .                                           | 27        |
| 6.8      | Dead Ends and Negative Results (Iteration 5) . . . . .                         | 27        |
| 6.9      | What Was Learned . . . . .                                                     | 28        |
| <b>7</b> | <b>Run Bookkeeping</b>                                                         | <b>28</b> |

|          |                                                            |           |
|----------|------------------------------------------------------------|-----------|
| <b>8</b> | <b>Hashes</b>                                              | <b>29</b> |
| <b>9</b> | <b>What the Run Learned Overall and What Is Still Open</b> | <b>29</b> |
| 9.1      | Summary of Confirmed Findings . . . . .                    | 29        |
| 9.2      | Dead Ends and Null Results . . . . .                       | 30        |
| 9.3      | Open Questions . . . . .                                   | 30        |

# 1 Introduction

This report documents an investigation of how emerging scientific concepts can be identified and characterised through evolving knowledge networks, and through which structural pathways they spread across disciplinary boundaries. The study builds on an OpenAlex-based scholarly dataset of 426 semantically grounded concepts with 462,812 works, constructs yearly co-occurrence and concept-discipline networks, and tests whether structural indicators provide early signals of emergence and cross-disciplinary integration.

Two research questions structure the work. RQ1 asks which temporal and structural patterns in an evolving, semantically grounded scientific knowledge graph characterise the emergence of scientific concepts. RQ2 asks how emerging scientific concepts diffuse across disciplinary communities over time, and which temporal network patterns distinguish locally concentrated concepts from concepts that become broadly integrated into the scientific knowledge network.

The study began with a source-sink viability hypothesis adapted from population ecology, asking whether a concept’s presence in a subfield is self-sustaining or import-dependent. That hypothesis was blocked by insufficient citation-layer density (Gate A failure: only 17.3% of edges meet the within-host tracing threshold). The investigation followed the pre-registered fallback.

For the emergence question (RQ1), the primary closure measure did not survive held-out confirmation (R1\_DEAD under the frozen kill rule: Holm  $p = 0.147$ ). Structural precursors of sustained uptake are volume and churn correlates. Only persistent-neighbour closure survives (Holm  $p = 0.015$ ).

For the diffusion question (RQ2), when a concept enters a new host subfield, the share of its initial co-occurrence partners that are native to that host predicts subsequent uptake by newcomer authors (held-out co-primary IRR/SD 1.19 [1.06, 1.33],  $p = 0.004$ ; MeSH replication IRR/SD 1.23 [1.12, 1.36], Holm  $p < 0.001$ ), while the share of origin companions does not.

The target venue is Applied Network Science.

## 2 Iteration 1: Data Foundation and Pre-Registration

### 2.1 Strategy

The first iteration was devoted to three preparatory tasks: (a) assembling a prior-art and pre-registration dossier grading each candidate mechanism’s novelty and feasibility, (b) building the concept pool and its associated OpenAlex work corpus, and (c) constructing the semantic grounding infrastructure (concept detection, labelling, variant merging) and a held-out MeSH confirmation population. No experiment or evaluation was executed. Every artifact in this iteration is either a dataset or a research dossier. The rationale was to fix all operational definitions, sample the data, and verify the citation-layer gate (Gate A) before committing compute to network construction and statistical modelling.

The pre-registration dossier ranks five candidate mechanisms by novelty margin against the nearest prior work. Host anchoring was ranked highest at medium-to-high novelty, toolkit co-transfer was ranked high on measurement novelty but coupled to host anchoring, citation-lineage source-sink viability was ranked medium, origin-neighbourhood structure was ranked low-to-medium, and the demic-versus-cultural adoption channel was ranked medium but with the weakest feasibility due to author-disambiguation biases.

A pre-registered screen was fixed before any data were inspected. The unit of analysis is a concept  $c$  in non-origin subfield  $d$  at focal year  $t$  (2010–2014), with at least 10  $d$ -papers in  $W_1 = [t-5, t]$ . The primary measure is out-of-sample ROC-AUC for establishment in  $W_2 = [t+1, t+5]$ ,

where establishment means at least 5  $W_2$   $d$ -papers by author-disjoint newcomers and presence in at least 3 of the 5  $W_2$  years.

## 2.2 Artifact 1: Concept Pool and OpenAlex Work Corpus

The concept pool was built in an outcome-blind frame. The sample was frozen and hashed before any post-appearance-window data were inspected (hash listed in Section 8). The pool comprises 206 concepts, of which 184 are main emerging concepts (first appearance year  $F$  in 2005–2016, with 20–300 papers in the window  $F$  to  $F+2$  and no more than 8,000 total works through 2024) and 22 stationary reference concepts (at least 5 hits in every year 1998–2004 with a 7-year sum of 70–700). The 184 main concepts are split by a SHA-1 hash into a screen fold of 123 and a held-out concept fold of 61, focal-year tags are separate (screen 2010–2014, held-out 2016–2018).

The pool is heavily dominated by arXiv-sourced concepts. Of 366 eligible candidates, 363 came from the arXiv-mined arm, and the final 184 span three origin fields.

Table 1: Concept pool composition by origin field (Iteration 1, `gen_art_dataset_1`).

| Origin field              | Main concepts | Share       |
|---------------------------|---------------|-------------|
| Physics and Astronomy     | 82            | 44.6%       |
| Physical Sciences (other) | 56            | 30.4%       |
| Computer Science          | 45            | 24.5%       |
| Other                     | 1             | 0.5%        |
| <b>Total main</b>         | <b>184</b>    | <b>100%</b> |

Table 2: Concept pool by first-appearance band (Iteration 1, `gen_art_dataset_1`).

| $F$ band     | Count      |
|--------------|------------|
| 2005–2007    | 38         |
| 2008–2011    | 76         |
| 2012–2016    | 70         |
| <b>Total</b> | <b>184</b> |

Verified links pairing each concept with the OpenAlex works that use it total 214,798 pairs spanning 208,374 unique works. Among the 121,829 main-arm links, 56,997 matched in the work title, 51,966 in the abstract, 10,177 through Semantic Scholar only, and 2,689 through OpenAlex concept indexing alone. [Correction, Iteration 2: the match-evidence counts above are main-arm only; across all 214,798 links the counts are `oa_abstract` 98,192, `oa_title` 92,043, `s2_only` 21,509, `oa_index_only` 3,054.] Of the 184 main concepts, 24 were discovered natively in OpenAlex and 160 were discovered through the Semantic Scholar index and mapped to OpenAlex (mapping rate 0.93). Twenty-one concepts carry a sense-check failure flag.

## 2.3 Gate A: Citation-Layer Feasibility

The quality report shows that among main-arm concept papers with references, 66.99% cite at least one earlier concept paper (the “traced share”), and after excluding first-year canonical papers and the five most-cited concept papers, the traced share remains 60.80%. The host-specific traced share

is 55.18%, above the 40% gate threshold on the lenient definition. [Correction, Iteration 2: the original description stated “at least one earlier concept paper in the same host subfield,” but the code counts any earlier concept paper as a parent regardless of subfield. The within-host share, computed in Iteration 2, is much lower (mean 0.20 on 724 edges), which is the finding that blocks the viability decomposition.]

Table 3: Citation-layer quality by origin field (Iteration 1, `gen_art_dataset_1`).

| Metric                           | All main | Phys. & Astro | Physical Sci | Computer Sci |
|----------------------------------|----------|---------------|--------------|--------------|
| Concept papers ( $n$ )           | 121,829  | 19,653        | 46,575       | 55,274       |
| Share with references            | 83.92%   | 82.04%        | 84.99%       | 83.74%       |
| Traced share (any parent)        | 66.99%   | 66.12%        | 67.89%       | 66.53%       |
| Traced share (excl. $F$ + top-5) | 60.80%   | 57.74%        | 61.75%       | 61.09%       |
| Host traced share                | 55.18%   | 49.68%        | 58.39%       | 53.86%       |
| Host concept papers ( $n$ )      | 41,861   | 4,171         | 15,973       | 21,609       |

The recall audit, comparing counts derived from OpenAlex against independent Semantic Scholar counts for 30 concepts, gives a median count ratio of 0.84 (IQR 0.79–0.89) and a Spearman rank correlation of 0.976.

## 2.4 Denominators and Venue Habitat

Totals by subfield and year cover 25,195 rows (all 252 OpenAlex subfields across 25 years in four counting variants). A venue-habitat file assigns 15,261 venues to their dominant subfield. [Correction, Iteration 2: the habitat was computed from each source’s all-years OpenAlex topic counts in the 2026 snapshot and is therefore neither pre-period nor citation-independent. An ASJC journal-level habitat was built in Iteration 2 as a replacement, though its coverage is only 12.1% of concept-work links.]

## 2.5 Artifact 2: Whole-Science Background Sample

A design-weighted background sample of 259,716 OpenAlex works from 1,260 strata (150 per field  $\times$  year, weights summing to 149,116,575) was built to supply co-occurrence network snapshots and calibrate host-nativeness profiles. It delivered 6,325 exact subfield  $\times$  year totals and 400 exact concept  $\times$  block subfield profiles for 100 calibration concepts at a cost of \$0.1952 in OpenAlex credits. [Correction, Iteration 5: this artifact is the workspace `gen_art_dataset_2`.]

The key finding of this artifact is negative: sample-based per-concept subfield profiles are unreliable for determining host-nativeness. In the calibration table, the median total-variation distance between sample and exact profiles is 0.41–0.79 by frequency decile, and top-subfield agreement is only 0.28–0.65.

The worker timed out after writing all outputs (no new JSONL records for 2,096 seconds), so its status is “failed” while its data are usable.

## 2.6 Artifact 3: Held-Out MeSH Confirmation Population

A separate biomedical confirmation arm was built from new MeSH descriptors (DateEstablished 2006–2016, widened to 2004–2005 and 2017–2018). Starting from 28,472 descriptors, the pipeline selects 5,201 topical descriptors, retains 3,430 after the provenance filter, narrows to 283 passing

the PubMed novelty pre-screen and early-volume rule, and adds widening-pool candidates. After the final rule, 191 concepts survive: 99 core 2006–2016, 14 from the 2004–2005 widening, and 78 from the 2017–2018 widening.

Table 4: MeSH held-out population by branch group (Iteration 1, gen\_art\_dataset\_3).

| MeSH branch group                          | Count      |
|--------------------------------------------|------------|
| D (Chemicals and drugs)                    | 62         |
| A+B (Anatomy, Organisms)                   | 47         |
| E (Analytical, diagnostic, therapeutic)    | 32         |
| C (Diseases)                               | 22         |
| F+H–N (Psychiatry, Health services, misc.) | 17         |
| G (Phenomena and processes)                | 11         |
| <b>Total</b>                               | <b>191</b> |

Caveats: only approximately 25% of new MeSH descriptors are genuinely new concepts; only 13 of 191 have their full non-PubMed works paged.

## 2.7 Artifact 4: Semantic Grounding and Labelling Infrastructure

The semantic grounding artifact comprises seven datasets. A stratified corpus of 93,600 English OpenAlex articles and reviews with abstracts was sampled at 150 per field  $\times$  year stratum across 26 OpenAlex fields. From 1.22 million distinct noun-phrase and acronym keys, 4,599 were labelled as CONCEPT, NOT\_CONCEPT, TOO\_GENERIC, or VARIANT\_OF by two LLMs with a claude-haiku-4.5 adjudicator.

Table 5: Labelling quality (Iteration 1, gen\_art\_dataset\_4).

| Metric                                            | Value |
|---------------------------------------------------|-------|
| Model A vs model B inter-rater kappa (4-class)    | 0.259 |
| Model A vs model B binary kappa (concept vs rest) | 0.267 |
| A vs adjudicator kappa (4-class, silver-gold 200) | 0.632 |
| A vs human anchors binary kappa (SemEval/SciERC)  | 0.56  |
| A precision (concept, human-anchored)             | 0.79  |
| A recall (concept, human-anchored)                | 0.76  |

The survivorship-free phrase pool yielded only 1 strict-eligible anchored phrase, far below the target of 150.

## 2.8 Artifact 5: Prior-Art Dossier and Pre-Registration

The research dossier grades each candidate mechanism’s novelty margin against the nearest prior art.

## 2.9 Dead Ends and Shortfalls (Iteration 1)

1. **Concept pool size.** 184 of 400 target concepts realised. The binding constraint was the free-singleton retrieval rate combined with heavy-tailed concept sizes.

Table 6: Candidate mechanism novelty grades (Iteration 1, gen\_art\_research\_1).

| Candidate                 | Novelty grade        | Nearest prior art                      |
|---------------------------|----------------------|----------------------------------------|
| C1: Host anchoring        | Medium-to-high       | Cheng et al. 2023                      |
| C2: Toolkit co-transfer   | High (coupled to C1) | No quantitative test found             |
| C0: Source-sink viability | Medium               | Kuhn et al. 2014; Maillart et al. 2026 |
| C4: Demic vs cultural     | Medium               | Fontaine et al. 2024                   |
| C3: Origin-neighbourhood  | Low-to-medium        | Structural diversity literature        |

2. **Field coverage.** The pool draws overwhelmingly from physics, physical sciences, and computer science (363 of 366 eligible candidates from the arXiv-mined arm).
3. **Survivorship-free phrase pool yield.** 1 of 150 target strict-eligible anchored concepts.
4. **Labelling quality.** Silver labels only (no human check). Model A’s bias is moderate; model B over-labels CONCEPT.

## 2.10 What Was Learned

Iteration 1 established the data foundation and pre-registration. The concept pool of 184 main emerging concepts and 22 reference concepts, with 214,798 verified concept-to-work links and 208,374 unique works, passes the citation-layer gate on the lenient (any-parent) definition but has not been tested at the within-host edge level required for viability estimation. No experiment was executed.

## 3 Iteration 2: Network Construction, Gate A Failure, and the Closure Lead

### 3.1 Strategy

Iteration 2 was a FIX iteration with five goals: (i) complete the hydration and add exact nativeness profiles and a citation-independent venue habitat; (ii) train and evaluate the concept grounding pipeline; (iii) compute Gate A as written (within-host, non-canonical, per  $W_1$  year) and estimate viability states with synthetic validation and pre-outcome power; (iv) build the yearly concept-concept co-occurrence network and run the RQ1 event study; (v) replicate the RQ1 event study on the MeSH population.

Pre-fixed decision rules for Iteration 3: (a) if Gate A fails, H1/H2 run with graft labels from the exact nativeness profiles; (b) H1 is declared a pilot if fewer than  $\sim 150$  concepts carry a tested edge; (c) RQ1 counts as confirmed only if at least one pre-named precursor has an event-study CI excluding 0 AND a delta-AUC CI excluding 0 on the main screen fold, keeps its sign on MeSH, and holds on the sealed held-out fold; (d) the held-out fold, the 2016–2018 focal years and the phrase pool remain untouched.

Note: the Iteration 2 RQ1 and viability experiments (Artifacts 8–10) ran on the Iteration 1 184-concept corpus in parallel with the hydration. The 426-concept hydrated corpus was not used until Iteration 3.

### 3.2 Artifact 6: Hydrated Corpus, Nativeness Profiles, and Venue Habitat

This artifact completes the hydration of the concept pool from 184 to 426 concepts (366 main: screen 247, held-out 119, 60 reference), with 462,812 works and 488,078 verified concept-work links. The frozen concept-frame hash was verified (see Section 8).

Table 7: Dataset 6 summary (Iteration 2, `gen_art_dataset` 5).

| Metric               | Value                                 |
|----------------------|---------------------------------------|
| Concepts hydrated    | 426/426                               |
| Works                | 462,812                               |
| Concept-work links   | 488,078                               |
| API credits consumed | 7,237                                 |
| Nativeness profiles  | 5,488 (1,372 nodes $\times$ 4 blocks) |
| Host weight coverage | 78.7%                                 |
| ASJC venues covered  | 2,929 (12.1% of links)                |

### 3.3 Artifact 7: Concept Grounding Pipeline

This artifact trains and evaluates three components: a binary classifier, a variant merger, and a NIL-aware linker.

Table 8: Classifier performance on test set ( $n = 300$ ) (Iteration 2, `gen_art_experiment` 1).

| Metric    | Classifier | LLM-A | Majority (predict-all-positive) |
|-----------|------------|-------|---------------------------------|
| F1        | 0.818      | 0.901 | 0.715                           |
| AUC       | 0.888      | —     | —                               |
| Precision | 0.759      | —     | —                               |
| Recall    | 0.886      | —     | —                               |
| Accuracy  | 0.780      | —     | —                               |

[Correction: previous values of P 0.812, R 0.824, accuracy 0.843 are not reproducible from the test predictions. The majority-class baseline F1 is 0.715, not 0.000 as previously stated. The classifier’s human-anchor kappa is 0.413; the 0.56 previously reported is LLM-A’s value.]

Against human anchors (SemEval-2017/SciERC): classifier E1 kappa 0.41, F1 0.674. E2 F1 0.706, AUC 0.82.

The variant merger operating at threshold  $p_{\text{merge}} = 0.855$  gives D3 test F1 0.357, held-out MeSH F1 0.222, B-cubed F1 0.78. The NIL-aware linker uses MiniLM embeddings,  $\tau = 0.95$  (link precision  $\geq 0.90$  at NIL prior 0.9, in-KB recall 0.404). Test population (hash in Section 8): MAIN 156, STRICT 147, REFERENCE\_ACCEPTED 42, SENSITIVITY 426.

### 3.4 Artifact 8: Viability Layer, Gate A, Viability States

All rules were pre-registered and hashed before any statistic (see Section 8).

### 3.4.1 Gate A Re-Test at Edge Level

**Gate A FAILS.** Only 17.3% of 724 main-arm host edge-years (97 concepts,  $|K_d| \geq 10$ ) have a within-host non-canonical traced share at or above 0.40. The mean within-host traced share is 0.20, median 0.15. On the same 724 edges, the lenient any-parent share is 0.477 (64.2%  $\geq 0.40$ ).




Figure 1: Gate A (citation-layer feasibility) on the same 724 concept-subfield-year host edges (97 concepts), comparing within-host citation tracing (blue; the Gate A definition) with lenient any-parent tracing (grey). Only 17.3% of edges (125/724) reach the threshold under within-host tracing, against 64.2% (465/724) under any-parent tracing. The within-host mean is 0.204 (median 0.15), against 0.477 for any-parent tracing. Because the drop comes from the tracing definition, the citation layer is too sparse for the viability decomposition, and Gate A fails and triggers the graft fallback route.

Table 9: Gate A edge-level results (Iteration 2, gen\_art\_experiment\_2).

| Metric                   | Within-host | Any-parent (same edges) |
|--------------------------|-------------|-------------------------|
| Share $\geq 0.40$        | 17.3%       | 64.2%                   |
| Mean traced share        | 0.20        | 0.477                   |
| Median traced share      | 0.15        | —                       |
| Reference arm host share | 0.13        | —                       |

Per pre-registered decision rule (a): Gate A  $0.1727 < 0.5$  triggers the graft fallback. The graft route has 2,347 host-entry events total, 2,154 with  $\geq 5$  entry-year keywords (screen 1,285, held-out 869).

### 3.4.2 Viability States (Descriptive Only)

Of 779 eligible edges, 169 were tested: SOURCE 63, SINK 26, FADING 7, UNDETERMINED 683 (reason codes: below\_nmin 65.6%, no\_cohort\_parents 12.7%, tested 21.7%). The synthetic FDR is 0.209, exceeding the 0.15 threshold. SOURCE FDR 0.13 (acceptable), SINK FDR 0.40 (unreliable).

### 3.4.3 Power Before Any Outcome

Per decision rule (b), H1 is PILOT-ONLY:  $N_c = 38$  at  $n_{\min} = 30$ , projected 77.5 (66–92) on the hydrated pool. Gate B fails: 0 labelled episodes.

## 3.5 Artifact 9: Precursor Event Study, Main Pool

This artifact runs the precursor event study on the main pool’s screen fold, testing whether volume-normalised structural precursors distinguish emerging from non-emerging concepts before onset. Built 25 yearly 3-year-window co-word snapshots (2000–2024) from the design-weighted background sample ( $\approx 27,000$  nodes,  $\approx 84,000$  kept edges, association-strength weights, Leiden best-of-5, alluvial IDs).

### 3.5.1 Emergence Labels

The primary label  $E$  (all-node strength percentile gain  $\geq 20$ ) yields only 4 onsets (rate 2.1%) because pool concepts sit in the bottom  $\sim 5\%$  of the legacy-concept strength distribution.  $E_{\text{alt}}$  (pool-only percentile, promoted post hoc): 14 onsets.  $E_{\text{up}}$  (sustained uptake only, also promoted post hoc): 41 onsets, 21 matched.

Table 10: Emergence label counts (Iteration 2, gen\_art\_experiment\_3).

| Label                          | $N$ emerging | Matched | Status                                 |
|--------------------------------|--------------|---------|----------------------------------------|
| $E$ (primary)                  | 4            | 3       | Underpowered pilot                     |
| $E_{\text{alt}}$ (sensitivity) | 14           | 10      | Post-hoc promotion                     |
| $E_{\text{up}}$ (uptake only)  | 41           | 21      | Post-hoc promotion, pilot ( $n < 30$ ) |

### 3.5.2 Event Study Results

Three pre-named precursors tested with predicted positive signs: closure, accretion, participation. All three pre-registered signs failed.

Table 11: Event study results,  $E_{\text{up}}$  (Iteration 2, gen\_art\_experiment\_3).

| Indicator            | Estimator    | $S$    | 95% CI          | Holm $p$ |
|----------------------|--------------|--------|-----------------|----------|
| Closure (log-ratio)  | Event study  | −0.837 | [−1.26, −0.427] | 0.003    |
| Closure (log-ratio)  | Pooled panel | −0.743 | [−1.06, −0.46]  | –        |
| Accretion share      | Event study  | −0.14  | [−0.185, 0.018] | 0.18     |
| Participation coeff. | Event study  | −0.007 | [−0.066, 0.056] | 0.824    |

[Correction: previous tables paired event-study  $S$  values with pooled-panel Holm  $p$ -values. The values above are estimator-consistent.]

Exploratory findings (BH  $q < 0.05$ ): before sustained uptake, concepts show higher novelty (+0.156), new-relation rate (+0.128), within-module  $z$ -score (+0.073), and burst state (+0.171). The artifact’s characterisation of emergence is “open, novel, fast-renewing neighbourhoods.”

### 3.5.3 Prediction

The pre-registered primary model is L2 logit. For  $E_{\text{up}}$ : BASELINE AUC 0.890, FULL (baseline + precursors) AUC 0.879,  $\Delta\text{AUC}$   $-0.010$   $[-0.045, 0.021]$ . Precursors add no predictive value beyond frequency, burst, degree, centrality and entropy baselines. [Correction: previous tables silently reported the secondary HistGradientBoosting model instead of the primary logit.]

### 3.5.4 Trajectory Patterns and Typology

Early bridging is common but non-discriminating (76% of all concepts). Incubation-then-expansion: 0% (concepts are “born expanding”). The 7-channel DTW typology is unstable at every  $k$  (min Jaccard  $\leq 0.51$ ). An entropy-only typology is stable (Jaccard 0.95, AMI 0.14 vs the network typology).

## 3.6 Artifact 10: Precursor Replication, MeSH Population

This artifact replicates the precursor event study on the 191 MeSH concepts.

Table 12: MeSH closure effects (Iteration 2, `gen_art_experiment_4`).

| Label           | $n_{\text{eff}}$ | Closure $D$ | 95% CI             | Holm $p$ |
|-----------------|------------------|-------------|--------------------|----------|
| PRIMARY         | 15               | $-0.086$    | $[-0.511, 0.402]$  | 0.677    |
| SENS1           | 29               | $-0.419$    | $[-0.740, -0.115]$ | 0.039    |
| $E_{\text{up}}$ | 51               | $-0.185$    | $[-0.498, 0.108]$  | 0.21     |

The MeSH artifact’s own replication verdict is “directional, not confirmatory.” Per decision rule (c): NOT MET.

MeSH prediction: precursors also add nothing beyond baselines (grouped-CV  $\Delta\text{AUC}$  +0.002  $[-0.035, 0.035]$ ).

## 3.7 Decision-Rule Evaluation

Table 13: Decision-rule evaluation (Iteration 2).

| Rule                                    | Verdict   | Evidence                                                                             |
|-----------------------------------------|-----------|--------------------------------------------------------------------------------------|
| (a) Gate A $\rightarrow$ graft fallback | TRIGGERED | Within-host $0.1727 < 0.5$ ; 2,347 graft events                                      |
| (b) H1 pilot-only                       | YES       | $N_c = 38 < 150$                                                                     |
| (c) RQ1 confirmed                       | NOT MET   | Sign not +; $\Delta\text{AUC}$ CI includes 0; MeSH sign inconsistent                 |
| (d) Held-out sealed                     | YES       | Exp. 3/Exp. 4 sealed; caveat: Exp. 2 computed origin series for 61 held-out concepts |

### 3.8 Dead Ends and Negative Results (Iteration 2)

1. **Gate A failure at edge level.** Within-host traced share 0.204 vs any-parent 0.477 on the same 724 edges. The viability decomposition cannot be executed.
2. **Synthetic validation FDR exceeds threshold.** Overall 0.209 > 0.15, SINK FDR 0.40.
3. **No stable 7-channel typology.** Jaccard  $\leq 0.51$  at every  $k$ .
4. **Precursors add no predictive value for WHETHER a concept emerges.**  $E_{\text{up}}$  logit  $\Delta\text{AUC}$   $-0.010$   $[-0.045, 0.021]$ .
5. **Closure is lower, not higher, before emergence.** The pre-registered + sign failed.

### 3.9 What Was Learned

Iteration 2 confirmed two dead ends (viability decomposition, prediction of emergence) and produced one lead (lower closure before onset). The lead is significant in one post-hoc main-pool label ( $E_{\text{up}}$ ,  $n = 21$ ,  $S = -0.837$ , Holm  $p = 0.003$ ), directionally supported in MeSH SENS1 ( $D = -0.419$ , Holm  $p = 0.039$ ), but with heterogeneous sign across labels.

## 4 Iteration 3: Deepening the Closure Lead and the Grafting Test

### 4.1 Strategy

Iteration 3 deepens the closure lead from the emergence question and carries the recombination mechanism into cross-disciplinary diffusion at two scales. The strategy (“Is it brokerage, and does grafting make it stick?”) has four components: (i) an evaluation artifact that audits every number in the Iteration 2 report; (ii) a deeper test of the emergence question on the hydrated 247-concept screen fold with turnover-proof openness measures; (iii) a breadth-prediction screen; (iv) a host-entry grafting test; and (v) a descriptive diffusion experiment.

### 4.2 Artifact 11: Record Audit and Turnover Pre-Check

This artifact audited 406 rows, 182 recomputed from item-level files. There are 42 non-OK drift flags: 14 WRONG\_ESTIMATOR, 10 WRONG\_UNITS, 9 WRONG\_DEFINITION, 8 DRIFT\_VALUE, 1 NOT\_REDERTIVABLE. These corrections have been incorporated into the Iteration 1 and Iteration 2 text.

Table 14: Turnover pre-check: residualised closure effect (Iteration 3, gen\_art\_evaluation\_1).

| Population               | $S_{\text{raw}}$ |        | $S_{\text{res}}$  | Retained | Verdict   |
|--------------------------|------------------|--------|-------------------|----------|-----------|
| Main ( $E_{\text{up}}$ ) | -0.758           | -0.566 | $[-1.140, 0.087]$ | 0.75     | Ambiguous |
| MeSH (SENS1)             | -                | -0.094 | $[-0.387, 0.217]$ | 0.60     | Ambiguous |
| IVW pooled               | -                | -0.186 | $[-0.453, 0.081]$ | -        | Ambiguous |

### 4.3 Artifact 12: Openness Before Take-Off

This artifact re-attaches all 426 hydrated concepts to the Iteration 2 co-word snapshots with vendored byte-identical code. Reproduction gate passed: code-mode closure  $r = 1.000000$  (max  $|\text{diff}| 5 \times 10^{-8}$ ).

Populations: after applying the frozen replacement rule (180 sense values filled): MAIN screen 202 (102 old, 100 new), held-out 100, STRICT 196/96, SENS 247/119.




Figure 2: Turnover-proof openness measures on the screen fold (MAIN population,  $E_{\text{up}}$  label; 68 onsets, 26 matched). Bars give the event-study effect size  $S$  for each measure, with whiskers showing 95% confidence intervals; the dashed vertical line marks  $S = 0$ . Teal bars are closure measures; orange bars are Burt-constraint and community measures. Burt constraint is positive, the opposite of the sign brokerage predicts.

Table 15: Emergence deepen results, MAIN  $\times$   $E_{\text{up}}$  (68 onsets, 26 matched) (Iteration 3, gen\_art\_experiment\_5).

| Measure                     | $S$    | 95% CI             | Interpretation                         |
|-----------------------------|--------|--------------------|----------------------------------------|
| Raw closure                 | -0.439 | $[-0.810, -0.050]$ | Lower closure before onset             |
| R1a residualised closure    | -0.295 | $[-0.663, 0.092]$  | $\sim 2/3$ retained; not pure turnover |
| R1b persistent-nbr closure  | -0.575 | $[-1.270, 0.122]$  | Same direction, underpowered           |
| R1c Burt constraint         | +0.083 | $[0.017, 0.147]$   | OPPOSITE to brokerage sign             |
| R1c cross-comm. pair excess | +0.063 | $[0.012, 0.111]$   | More cross-community ties              |
| R1d hub-not-clique          | +0.061 | $[-0.009, 0.133]$  | Not significant                        |

Mechanical verdict: MIXED. The closure effect is not reducible to neighbourhood turnover

(residualised closure retains  $\sim 2/3$ ), but it is not Burt brokerage either: Burt constraint is higher, not lower, before onset.

Pooled panel (all 68 onsets, ungrouped): persistent-neighbour closure  $-1.12$   $[-1.64, -0.71]$  (Holm  $p = 0.0015$ ).

Fresh replication on 100 new concepts: raw  $S = -0.455$   $[-1.006, 0.060]$ , one-sided  $p = 0.044$ .

Prediction: grouped CV  $\Delta\text{AUC} -0.001$   $[-0.008, 0.006]$ ; precursors do not predict WHETHER a concept will show sustained uptake.

#### 4.4 Artifact 13: Breadth-Prediction Screen

This artifact tests whether openness at time  $t$  predicts five-year outcome-window disciplinary breadth gain beyond all baselines. Population: MAIN screen 202 concepts, held-out 100 sealed.

Table 16: Breadth-prediction screen results (Iteration 3, gen\_art\_experiment\_6).

| Openness measure | Outcome                              | $\beta/\text{SD}$ | 95% CI            | Holm $p$ |
|------------------|--------------------------------------|-------------------|-------------------|----------|
| closure_res      | $Y_{1r}$ (Shannon change)            | +0.019            | $[-0.010, 0.041]$ | 0.405    |
| closure_res      | $Y_2$ (Rao-Stirling change)          | +0.007            | $[-0.002, 0.013]$ | 0.405    |
| closure_res      | $\log(1 + Y_3)$ (newcomer subfields) | -0.065            | $[-0.156, 0.030]$ | 0.405    |
| closure_res_imp  | $Y_2$ (co-primary)                   | +0.012            | $[0.002, 0.021]$  | 0.06     |

Result: breadth prediction not supported on the screen fold. Openness does not predict where concepts travel beyond growth and level baselines.

#### 4.5 Artifact 14: Host-Entry Grafting Test

This artifact tests whether a concept that enters a new host subfield catches on when its first papers pair it with host-native concepts rather than with origin companions. Pre-registration frozen before any outcome (see Section 8).

Setup: 4,177 main-arm host entries (concept  $c$  first appears in non-origin subfield  $d$  in year  $e$ ). Primary sample: 1,740 screen MAIN entries with  $\geq 5$  partners, in 184 concepts. Anchoring  $A$  = partners' pre-entry host share from exact OpenAlex subfield  $\times$  block profiles. Co-transfer  $CT$  = share of partners that were origin companions in  $[e-5, e-1]$ . Entry is mostly a package: 70% non-native origin companions, 1.8% native grafts.

Table 17: Host-entry grafting results (Iteration 3, gen\_art\_experiment\_7).

| Specification                                   | $A$ (IRR/SD) | 95% CI         | Holm $p$           | $CT$ ( $p$ ) | Verdict  |
|-------------------------------------------------|--------------|----------------|--------------------|--------------|----------|
| Primary (concept $\times e$ + host $\times e$ ) | 1.39         | $[0.97, 1.99]$ | 0.15               | 0.85         | NEITHER  |
| Co-primary (concept + $e$ + host)               | 1.30         | $[1.16, 1.45]$ | $1 \times 10^{-5}$ | 0.48         | GRAFTING |
| Concept + host $\times e$                       | 1.22         | —              | —                  | —            | GRAFTING |
| Concept $\times$ 2-yr bin                       | —            | —              | —                  | —            | NEITHER  |

[Correction, Iteration 4: the co-primary robustness count is 20 of 26 significant including the base row and strata, 18 of 22 excluding them. The six null rows are: binary native  $\geq 0.5$  (IRR 1.005,  $p = 0.92$ ), binary native  $\geq 0.7$  (1.066,  $p = 0.29$ ),  $A_{\text{distinct}}$  (1.003,  $p = 0.96$ ), exclude

single-paper entries (1.244 [0.90, 1.71],  $p = 0.17$ ,  $N = 200$ ,  $G = 50$ ), CS stratum (1.15,  $p = 0.11$ ), and Physics/Astronomy stratum (0.99,  $p = 0.95$ ).]

## 4.6 Artifact 15: Descriptive Diffusion

### 4.6.1 Diffusion Typology

3-channel diffusion typology (rarefied Shannon, Rao-Stirling, active subfields, normalised multivariate DTW +  $k$ -medoids, Hennig bootstrap  $B = 200$ ). Only  $k = 2$  is stable (min Jaccard 0.861): **localised** ( $n = 136$ ) and **broad from the start** ( $n = 66$ ).

Table 18: Diffusion typology summary (Iteration 3, gen\_art\_experiment\_8).

| Cluster          | $n$ | Description                                      | Dominant origin |
|------------------|-----|--------------------------------------------------|-----------------|
| Localised        | 136 | Low entropy, few active subfields                | Physics         |
| Broad from start | 66  | High entropy and many subfields from early years | Mixed           |

The entropy-only baseline typology is stable at finer  $k = 4$  (Jaccard 0.856). After residualising on early volume, the 3-channel typology does NOT separate unclustered outcomes better than the entropy-only baseline.

### 4.6.2 Community Roles

BRIDGE 0.61, OTHER 0.33, STAYER 0.03, MIGRANT 0.02; pre-declared CORE\_GROWING and FOUNDER roles never fire. Lagged roles do not robustly predict host-subfield entry: BRIDGE OR 0.64 [0.38, 1.08].

### 4.6.3 Expansion Precedes Diffusion

In 15 of 16 concepts exhibiting both an expansion onset and a diffusion onset, expansion precedes diffusion (proportion 0.94 [0.81, 1.0], year-shuffle null 0.62,  $p = 0.004$ ). Underpowered (16 concepts).

## 4.7 Dead Ends and Negative Results (Iteration 3)

1. **Openness does not predict WHERE concepts travel.** Breadth-prediction screen: all Holm  $p = 0.405$ .
2. **Brokerage is not the mechanism behind low closure.** Burt constraint is higher, not lower, before onset ( $S = +0.083$  [0.017, 0.147]).
3. **The 3-channel diffusion typology adds nothing beyond entropy.** After residualising on volume, the multi-channel typology does not separate outcomes better than entropy-only.
4. **Community roles do not predict host entry.** BRIDGE OR 0.64 [0.38, 1.08].
5. **Co-transfer is null.** Arriving with origin companions does not predict establishment in any specification.




Figure 3: Representative cases of the  $k = 2$  diffusion typology: localised ( $n = 136$ , top row) and broad-from-the-start ( $n = 66$ , bottom row). Each row shows the cluster's medoid (left) and the nearest non-medoid concept (right). Localised concepts stay near-monopolised by a single subfield (88–97%, entropy 0.14–0.40 nats). Broad concepts split across two or more subfields from their first window (entropy 0.79–1.24 nats).

## 4.8 What Was Learned

Iteration 3 advances the study on three fronts. The closure effect survives turnover controls (residualised closure retains  $\sim 2/3$  of the raw effect) but is not Burt brokerage (constraint is higher, not lower). Host-entry grafting predicts durable integration on the screen fold (co-primary IRR 1.30 [1.16, 1.45], Holm  $p = 1 \times 10^{-5}$ ), holding across 20 of 26 robustness variants. Co-transfer is null. Breadth prediction is null.

## 5 Iteration 4: Confirmation on Held-Out Folds and Cross-Population Replication

### 5.1 Strategy

Iteration 4 is a CONFIRMATION iteration. Its purpose is to open the four sealed held-out folds and determine which screen-fold findings replicate, to run four discriminating tests (single-paper artefact, vocabulary decomposition, topic-score absorption, and MeSH replication) that separate genuine host-vocabulary effects from artefacts and circularity, to replicate the grafting result on the independent MeSH population, to test whether pre-emergence closure predicts later host-entry anchoring, and to test the adopter-level mechanism.

### 5.2 Reviewer Critique Disposition

Table 19: Reviewer critique disposition (Iteration 4).

| #   | Action taken                                                                             |
|-----|------------------------------------------------------------------------------------------|
| M1  | Chronology broken: iterations 1–2 silently rewritten → Inline [Correction] markers added |
| M2  | Iteration-3 tables missing detail rows → Rows incorporated                               |
| M3  | Grafting robustness count wrong (26/27 → 20/26) → Corrected inline                       |
| M4  | Closure claim overstated → Addressed by held-out confirmation below                      |
| M5  | Lead-lag 94% disproportionate → Full category-count table added                          |
| M6  | Iteration-3 reasoning partly recorded → Strategy expanded                                |
| M7  | Coverage table overstates completion → Updated with caveats                              |
| M8  | Traceability: folder names instead of artifact IDs → Corrected                           |
| M9  | Prior-review must-fix items still missing → Incorporated where data available            |
| M10 | Nearest-neighbour comparisons one-sided → Addressed with Cheng et al. tension            |
| m1  | Minor traceability and labelling slips → Corrected inline                                |

### 5.3 Artifact 16: Held-Out Grafting Confirmation and Discriminating Tests

The held-out and discriminating-test specification hashes were verified before any outcome was read (see Section 8).

#### 5.3.1 Held-Out Grafting Confirmation

The co-primary specification confirms the grafting effect on the held-out fold: IRR/SD = 1.19 [1.06, 1.33],  $p = 0.0045$ , wild-cluster  $p = 0.012$ . The primary specification is inconclusive because it retains only 23% of events (250 of 1,097) with 30 concept clusters.

Table 20: Held-out grafting results (Iteration 4, `gen_art_evaluation_2`).

| Spec       | Fold     | IRR/SD | 95% CI       | $p$       | CT   | $N$   | $G$ | Verdict      |
|------------|----------|--------|--------------|-----------|------|-------|-----|--------------|
| Primary    | Held-out | 0.98   | [0.70, 1.37] | 0.91      | 0.82 | 250   | 30  | Inconclusive |
| Co-primary | Held-out | 1.19   | [1.06, 1.33] | 0.0045    | 0.60 | 972   | 74  | CONFIRMED    |
| Co-primary | Screen   | 1.30   | [1.16, 1.45] | $10^{-5}$ | –    | 1,544 | 140 | GRAFTING     |




Figure 4: Host-entry grafting across folds. Each row shows the incidence-rate ratio for newcomer uptake per SD of  $A_{\text{cont}}$  (host-nativeness of entry partners), with whiskers marking its confidence interval. The co-primary effect replicates on the held-out fold (IRR/SD 1.19 [1.06, 1.33]) and the MeSH fold (1.23 [1.12, 1.36]), and pools to 1.26 [1.17, 1.36].

### 5.3.2 Single-Paper Artefact Test

87% of screen co-primary events (1,344 of 1,544) are single-paper entries.

Table 21: Single-paper artefact interaction results (Iteration 4, `gen_art_evaluation_2`).

| Fold     | IRR/SD (single) | IRR/SD (multi) | Ratio | $p$ (interaction) |
|----------|-----------------|----------------|-------|-------------------|
| Pooled   | 1.26            | 1.28           | 0.92  | 0.11              |
| Held-out | 1.37            | 1.19           | 0.82  | 0.004             |

### 5.3.3 Host-Vocabulary Decomposition

Table 22: Vocabulary decomposition, co-primary spec (Iteration 4, `gen_art_evaluation_2`).

| Component | Fold     | IRR/SD | 95% CI       | $p$                |
|-----------|----------|--------|--------------|--------------------|
| NATIVE    | Pooled   | 1.11   | [1.05, 1.19] | 0.0008             |
| ADJACENT  | Pooled   | 1.32   | [1.20, 1.46] | $9 \times 10^{-8}$ |
| NATIVE    | Held-out | 1.10   | [1.01, 1.20] | 0.036              |
| ADJACENT  | Held-out | 1.19   | [1.03, 1.38] | 0.020              |

Both NATIVE and ADJACENT components are positive on both folds. On held-out, neither is significant after Holm correction (NATIVE Holm  $p = 0.071$ , ADJACENT Holm  $p = 0.059$ ).

### 5.3.4 Topic-Score Absorption

Combined topic-score absorption:  $-7.3\%$   $[-27.5, +6.7]$  of the log-IRR on the pooled fold,  $-10.3\%$   $[-84.3, +56.6]$  on the held-out fold. The CI includes zero on both folds: topic-score controls do not meaningfully reduce the  $A_{\text{cont}}$  effect.

## 5.4 Artifact 17: MeSH Replication of Grafting

This artifact replicates the grafting test on the 191 MeSH biomedical concepts (spec hash in Section 8).

Table 23: MeSH grafting results (Iteration 4, `gen_art_experiment_9`).

| Spec           | IRR/SD | 95% CI       | Holm $p$             | CT IRR/SD | CT $p$ | $N$   | $G$ |
|----------------|--------|--------------|----------------------|-----------|--------|-------|-----|
| R1: primary    | 1.32   | [1.10, 1.60] | 0.007                | 1.10      | 0.46   | 1,004 | 122 |
| R2: co-primary | 1.23   | [1.12, 1.36] | $9.7 \times 10^{-5}$ | 1.10      | 0.053  | 2,171 | 160 |

**Verdict: REPLICATED.** Both the primary and co-primary specifications are significant on MeSH. The IVW pooled estimate across main-arm co-primary (IRR/SD 1.30) and MeSH co-primary (1.23) is 1.26 [1.17, 1.36] with  $I^2 = 0$  (no heterogeneity).

## Closure event study: screen vs. held-out fold




Figure 5: Pre-emergence closure effect sizes  $S$  (markers) with 95% confidence intervals (whiskers), for four closure measures on the screen fold (blue circles) and the held-out fold (vermillion diamonds). The primary raw-closure measure fails confirmation (R1\_DEAD). Persistent-neighbour closure stays negative on held-out data ( $S = -1.07$ , 95% CI  $[-1.86, -0.29]$ , Holm  $p = 0.015$ ; CONFIRMED). Burt constraint is positive on both folds, inconsistent with a brokerage explanation.

Table 24: Closure held-out results (Iteration 4, gen\_art\_evaluation\_3).

| Measure                 | Screen $S$ | Held-out $S$ | 95% CI             | Holm $p$ | Decision             |
|-------------------------|------------|--------------|--------------------|----------|----------------------|
| Closure (pooled panel)  | -0.661     | -0.391       | $[-0.855, 0.072]$  | 0.147    | DEAD                 |
| Turnover-resid. closure | -0.400     | -0.109       | $[-0.561, 0.342]$  | 0.635    | DEAD                 |
| Persistent-nbr closure  | -1.124     | -1.073       | $[-1.858, -0.289]$ | 0.015    | CONFIRMED            |
| Burt constraint         | +0.083     | +0.105       | $[0.021, 0.187]$   | -        | DEAD<br>(wrong dir.) |

## 5.5 Artifact 18: Closure Held-Out Confirmation

The held-out fold has 100 MAIN concepts (spec hash in Section 8). The matching procedure yields 30 onsets, 22 matched.

Overall closure status: not confirmed on the primary operationalisation. Only persistent-neighbour closure (restricted to top-20 neighbours present in consecutive years) is confirmed on the held-out fold.

### 5.5.1 Closure-Anchoring Link Test

Table 25: Closure-anchoring link test (Iteration 4, `gen_art_evaluation_3`).

| Fold     | $n$ | Partial $\rho$ | 95% CI            | $p$   |
|----------|-----|----------------|-------------------|-------|
| Screen   | 102 | +0.025         | $[-0.199, 0.242]$ | 0.806 |
| Held-out | 49  | +0.053         | $[-0.279, 0.391]$ | 0.736 |
| Pooled   | 151 | +0.049         | $[-0.126, 0.228]$ | 0.564 |

Verdict: NULL. There is no detectable association between pre-emergence closure and later host-entry anchoring.

## 5.6 Artifact 19: Descriptive Held-Out Confirmation

### 5.6.1 Lead-Lag: Expansion Before Diffusion

Table 26: Lead-lag category counts (Iteration 4, `gen_art_evaluation_4`).

| Category        | Screen ( $n=202$ ) | Held-out ( $n=100$ ) | Pooled main ( $n=302$ ) | MeSH ( $n=191$ ) |
|-----------------|--------------------|----------------------|-------------------------|------------------|
| Neither         | 110 (54%)          | 49 (49%)             | 159 (53%)               | 71 (37%)         |
| Diffusion only  | 39 (19%)           | 23 (23%)             | 62 (21%)                | 97 (51%)         |
| Expansion only  | 37 (18%)           | 18 (18%)             | 55 (18%)                | 10 (5%)          |
| Expansion first | 15 (7%)            | 10 (10%)             | 25 (8%)                 | 9 (5%)           |
| Same year       | 1 (0.5%)           | 0                    | 1 (0.3%)                | 1 (0.5%)         |
| Diffusion first | 0                  | 0                    | 0                       | 3 (1.6%)         |

Among concepts with both transitions observable, expansion precedes diffusion in 25 of 26 pooled main-arm concepts (held-out: 10 of 10, screen: 15 of 16). On MeSH, 9 of 13 dual-onset concepts show expansion first, with 3 showing diffusion first ( $p = 0.43$ : not different from the null).

### 5.6.2 Held-Out Confirmation of Descriptive Findings

Five pre-registered expectations were tested on the held-out fold: typology shares (PASS), type  $\times$  outcome association (PASS), expansion-first ordering (PASS), anchoring  $\times$  typology interaction (FAIL: lagged-role entry ORs flip sign), other patterns (PASS).

### 5.6.3 Rooting

BROAD concepts do not anchor more: screen mean  $A_{\text{cont}}$  difference  $-0.0018$ , held-out  $+0.0005$ , CI including zero.

Table 27: Rooting contrasts (Iteration 4, gen\_art\_evaluation\_4).

| Type      | Entries/concept | EST rate | Raw $A_{\text{cont}}$ | $CT$  | PPML IRR/SD | 95% CI       |
|-----------|-----------------|----------|-----------------------|-------|-------------|--------------|
| BROAD     | 18.4            | 0.31     | 0.039                 | 0.584 | 1.36        | [1.20, 1.55] |
| LOCALISED | 6.6             | 0.13     | 0.056                 | 0.691 | 1.16        | [0.98, 1.38] |

### 5.7 Artifact 20: Adopter-Level Mechanism

This artifact tests whether authors who adopt a concept in a new host subfield had prior corpus exposure to the concept’s entry partners.

Table 28: Adopter enrichment results (Iteration 4, gen\_art\_evaluation\_5).

| Exposure measure                                        | OR   | 95% CI (boot) | $p$ (CRV)             |
|---------------------------------------------------------|------|---------------|-----------------------|
| $E_{\text{any}}$ (any prior partner exposure)           | 3.09 | [2.32, 4.28]  | $1.4 \times 10^{-13}$ |
| $E_{\text{neg}}$ (frequency-matched negative controls)  | 0.62 | [0.49, 0.77]  | $2.6 \times 10^{-5}$  |
| $E_{\text{plac}}$ (partners of another concept’s entry) | 0.72 | [0.57, 0.95]  | 0.013                 |
| $E_{\text{swap}}$ (exposure swap control)               | 1.11 | [0.86, 1.39]  | 0.39                  |

Table 29: Vocabulary-class enrichment (Iteration 4, gen\_art\_evaluation\_5).

| Exposure class                               | OR   | 95% CI (boot) | $p$ (CRV)   |
|----------------------------------------------|------|---------------|-------------|
| $E_{\text{for}}$ (FOREIGN partner exposure)  | 2.88 | [2.28, 3.69]  | $< 10^{-6}$ |
| $E_{\text{adj}}$ (ADJACENT partner exposure) | 1.91 | [1.43, 2.73]  | $10^{-4}$   |
| $E_{\text{nat}}$ (NATIVE partner exposure)   | 1.47 | [1.02, 2.33]  | 0.059       |
| $E_{\text{neg}}$ (non-partner exposure)      | 0.64 | [0.51, 0.82]  | 0.0003      |

The gradient runs FOREIGN > ADJACENT > NATIVE. Adopters are most enriched for exposure to the concept’s foreign (origin-vocabulary) partners, consistent with absorptive capacity or topical proximity. This runs in the opposite direction from the entry-level grafting finding: at the entry level, host-native partners facilitate uptake by newcomers who do not need origin expertise; at the adopter level, the individuals who adopt are those with prior origin-vocabulary exposure.

### 5.8 Dead Ends and Negative Results (Iteration 4)

1. **Closure not confirmed on held-out for the primary measure.** Raw closure  $S = -0.391$ , Holm  $p = 0.147$ ; turnover-residualised closure  $S = -0.109$ , Holm  $p = 0.635$ . Only persistent-neighbour closure survives.
2. **Closure-anchoring link: null.** Pre-emergence closure does not predict later host-entry anchoring.
3. **BROAD concepts do not anchor more.** The hypothesis is not supported.
4. **Grafting primary specification: inconclusive.** IRR/SD 0.98 [0.70, 1.37] on held-out, 250 events / 30 clusters, declared underpowered in advance.

5. **Single-paper interaction on held-out.** Single-paper entries drive the effect more strongly on the held-out fold (ratio 0.80,  $p = 0.004$ ).
6. **Anchoring  $\times$  typology interaction flips between folds.** Not confirmed.
7. **Breadth prediction: all six sensitivity cells null on held-out.**

## 5.9 What Was Learned

Iteration 4 resolves the study’s two main questions with held-out and cross-population evidence.

The co-primary grafting specification replicates on the held-out fold (IRR/SD 1.19 [1.06, 1.33],  $p = 0.0045$ ) and on the independent MeSH population (IRR/SD 1.23 [1.12, 1.36], Holm  $p = 9.7 \times 10^{-5}$ ). The IVW pooled estimate is 1.26 [1.17, 1.36] with  $I^2 = 0$ . Co-transfer is null across all folds and populations.

The primary closure measure did not survive held-out confirmation (Holm  $p = 0.147$ ). Only persistent-neighbour closure is confirmed ( $S = -1.073$ , Holm  $p = 0.015$ ).

## 6 Iteration 5: Post-Confirmation Decomposition and Full Audit

### 6.1 Strategy

Iteration 5 is a POST-CONFIRMATION EXPLORATORY iteration. The grafting result was confirmed in Iteration 4 and replicated on MeSH. Three post-confirmation tests decompose the finding along dimensions the pre-registered confirmation did not distinguish: (i) extensive vs intensive margin, (ii) host-specific vs generic vocabulary, (iii) field boundary. A full numerical audit was run, along with a positioning dossier and publication-ready figure specifications.

### 6.2 Artifact 21: Margin Decomposition and Field-Boundary Test

Spec hash listed in Section 8.

#### 6.2.1 Extensive vs Intensive Margin

Table 30: Extensive-margin results (LPM, pp change per SD of  $A_{\text{cont}}$ ) (Iteration 5, gen\_art\_evaluation\_6).

| Fold     | pp/SD | 95% CI        | $p$ (CRV1)            | $N$   | $G$ |
|----------|-------|---------------|-----------------------|-------|-----|
| Screen   | +6.46 | [3.42, 9.51]  | $4.5 \times 10^{-5}$  | 1,686 | 169 |
| Held-out | +7.29 | [2.97, 11.61] | 0.0012                | 1,046 | 85  |
| MeSH     | +6.17 | [3.87, 8.46]  | $3.4 \times 10^{-7}$  | 2,236 | 176 |
| IVW      | +6.43 | [4.76, 8.11]  | $5.1 \times 10^{-14}$ | –     | –   |

Margin-decomposition verdict: BOTH. Host-leaning entry vocabulary raises both the probability that uptake starts (extensive) and its magnitude given it has started (intensive). The extensive share is 38% [21%, 55%] of the total log-IRR. The binary establishment threshold (EST\_bin) is null on the held-out fold (+2.54 pp,  $p = 0.16$ ).




Figure 6: Overview of the study design, drawn as four bands linked by arrows. (a) Concept pool: 426 concepts from OpenAlex, split into 366 main and 60 reference concepts. (b) Co-word network: 25 yearly co-occurrence snapshots ( $\approx 27,000$  nodes,  $\approx 84,000$  edges) partitioned into Leiden communities. (c) Analysis: RQ1 (emergence) returns R1\_DEAD on the primary closure test; persistent-neighbour closure is confirmed. RQ2 (diffusion) tests host-entry grafting, confirmed (IRR/SD = 1.19, held-out; 1.23, MeSH). K1 (BOTH margins) and K2 (HOST-SPECIFIC) are supported; K3 finds no field boundary. (d) Diffusion typology: a  $k = 2$  solution separating localised concepts (136) from concepts broad from the start (66).

Table 31: Intensive-margin results (PPML on  $Y \geq 1$  subsample, IRR/SD) (Iteration 5, `gen_art_evaluation_6`).

| Fold     | IRR/SD | 95% CI         | $p$ (CRV1)           | $N$   | $G$ |
|----------|--------|----------------|----------------------|-------|-----|
| Screen   | 1.272  | [1.121, 1.443] | 0.0003               | 795   | 107 |
| Held-out | 1.163  | [1.008, 1.342] | 0.038                | 489   | 58  |
| MeSH     | 1.137  | [1.026, 1.260] | 0.015                | 1,169 | 143 |
| IVW      | 1.183  | [1.104, 1.267] | $1.7 \times 10^{-6}$ | —     | —   |

Table 32: Field-boundary results (co-primary PPML, IRR/SD by origin field) (Iteration 5, `gen_art_evaluation_6`).

| Origin field      | IRR/SD | 95% CI         | $p$                  |
|-------------------|--------|----------------|----------------------|
| Physics/Astronomy | 1.178  | [0.926, 1.499] | 0.18                 |
| Computer Science  | 1.216  | [1.066, 1.386] | 0.004                |
| Other             | 1.281  | [1.167, 1.407] | $3.5 \times 10^{-6}$ |
| MeSH (reference)  | 1.233  | [1.120, 1.360] | $9.7 \times 10^{-5}$ |

### 6.2.2 Field Boundary

Wald test for equality:  $W = 0.900$ ,  $p = 0.638$  (CRV1), label-permutation  $p = 0.770$ . Field-boundary verdict: NO DETECTABLE FIELD BOUNDARY.

### 6.3 Artifact 22: Host-Specificity Test

Spec hash listed in Section 8.

Table 33: Host-specificity retention:  $A_{\text{cont}}$  coefficient stability after adding generality controls (Iteration 5, gen\_art\_evaluation\_7).

| Fold     | Retention | 95% CI         | % removed | $G_H$ IRR/SD | $G_F$ IRR/SD |
|----------|-----------|----------------|-----------|--------------|--------------|
| Screen   | 1.006     | [0.733, 1.312] | −0.6%     | 1.122        | 0.856        |
| Held-out | 1.062     | [0.438, 1.802] | −6.2%     | 1.068        | 0.943        |
| MeSH     | 0.876     | [0.640, 1.037] | +12.4%    | 1.110        | 0.764        |
| IVW      | 0.969     | [0.81, 1.10]   | +3.1%     | 1.103        | 0.815        |

Table 34:  $A_{\text{lift}}$  results (PPML, IRR/SD) (Iteration 5, gen\_art\_evaluation\_7).

| Fold     | Model                                | IRR/SD | 95% CI         | Holm $p$              |
|----------|--------------------------------------|--------|----------------|-----------------------|
| Screen   | M3 ( $A_{\text{lift}} + G_H + G_F$ ) | 1.661  | [1.361, 2.026] | $1.4 \times 10^{-6}$  |
| Held-out | M3                                   | 1.459  | [1.146, 1.856] | 0.003                 |
| MeSH     | M3                                   | 1.244  | [1.120, 1.381] | $6.3 \times 10^{-5}$  |
| IVW      | M3                                   | 1.340  | [1.232, 1.458] | $1.5 \times 10^{-10}$ |

Host-specificity verdict: HOST-SPECIFIC. Generality controls do not reduce the  $A_{\text{cont}}$  coefficient (IVW retention 0.969 [0.81, 1.10], 3.1% removed). The Balassa-index  $A_{\text{lift}}$  measure replicates across all folds (IVW IRR/SD 1.34 [1.23, 1.46]).

Caveat on  $A_{\text{lift}}$ : the within-FE correlation between  $A_{\text{lift}}$  and  $\ln(A_{\text{cont}})$  is 0.997, so the lift leg is weak in adding independent evidence.

### 6.4 Artifact 23: Full Numerical Audit

Table 35: Audit module results (Iteration 5, gen\_art\_evaluation\_8).

| Module | Description                  | Result                                    |
|--------|------------------------------|-------------------------------------------|
| M1     | Curated-number accuracy      | 99.4% (319 curated)                       |
| M2     | Headline + independent check | 90.8% (210 R-tier 92.9%, 50 C-tier 82.0%) |
| M5     | Seeded-perturbation recall   | 0.60 overall                              |
| M6     | Independent claim review     | 73/73 agree; 0 blocked claims             |
| M7     | Table-cell traceability      | 11/11 tables pass                         |
| M8     | Reviewer-critique closure    | All 6 MAJOR items CLOSED                  |
| M9     | Decision-rule consistency    | 13 rules; 1 mismatch                      |

Of 45 drift rows, 6 are VERDICT-CHANGING, 16 are WORDING, and 23 are NUMBER-ONLY.

### 6.4.1 Table 18 Self-Test Honesty

Table 36: Table 18 claims vs audit findings (Iteration 5, gen\_art\_evaluation\_8).

| Claim                          | Audit status | Note                                         |
|--------------------------------|--------------|----------------------------------------------|
| M1: [Correction] markers added | PARTIAL      | Drift lines 7, 50, 77, 477, 551 lack markers |
| M2: Holm columns added         | NOT DONE     | Table 14 header has no Holm column           |
| M3: robustness count corrected | PARTIAL      | “26 of 27” still appears at 3 locations      |
| M4: closure claim graded       | PARTIAL      | R1_DEAD not stated at summary                |
| M5: lead-lag table added       | DONE         | MeSH null $p = 0.43$ still missing           |
| M6: decision rules recorded    | NOT DONE     | Strategy is 5 lines                          |
| M7: coverage table updated     | DONE         | –                                            |
| M8: artifact IDs used          | PARTIAL      | Artifact 2 marker mismatch                   |
| M9: prior must-fix items       | NOT DONE     | 292 MISSING_IN_REPORT rows                   |
| M10: nearest-nbr two-sided     | DONE         | –                                            |
| m1: minor slips corrected      | NOT DONE     | –                                            |

## 6.5 Representative-Case Interpretations

Table 37: Per-case entry statistics, co-primary sample (Iteration 5, gen\_art\_evaluation\_8).

| Concept                 | Entries | Rooted | Mean $A_{\text{cont}}$ (rooted) | Mean $A_{\text{cont}}$ (unrooted) | Diff   |
|-------------------------|---------|--------|---------------------------------|-----------------------------------|--------|
| Wireless backhaul       | 14      | 1      | 0.084                           | 0.021                             | +0.062 |
| EPR steering            | 7       | 1      | 0.142                           | 0.027                             | +0.115 |
| Locally repairable code | 1       | 1      | 0.056                           | –                                 | –      |
| Holographic QCD         | 5       | 0      | –                               | 0.041                             | –      |

## 6.6 Artifact 24: Positioning Dossier

Novelty verdict: PARTLY ANTICIPATED (0 ANTICIPATES, 11 PARTLY among 25 graded candidates). No located work measures how host-specific the vocabulary is that a concept is combined with when it enters a new subfield. The nearest published result for the diffusion question is Cheng et al. (2023), who measure global ideational embeddedness of new ideas with a yearly article-count outcome (not a binary “core” status as previously implied), and find +1 SD  $\rightarrow$  approximately +25% articles.

## 6.7 Artifact 25: Paper Figures

All 8 figures pass production checks: correct dimensions (174 mm width), 600 DPI, no Type 3 fonts. Total: 957 values checked, 0 failures, value fidelity 1.0.

## 6.8 Dead Ends and Negative Results (Iteration 5)

1. **EST\_bin null on held-out.** The binary establishment threshold shows +2.54 pp per SD on held-out with  $p = 0.16$ . The extensive-margin result holds for uptake initiation ( $Y \geq 1$ ) but not for strict establishment.

Table 38: Figure inventory and production checks (Iteration 5, `gen_art_evaluation_9`).

| Figure | Description       | Values checked | All pass |
|--------|-------------------|----------------|----------|
| F1     | Method flow       | 39 exact       | Yes      |
| F2     | D2 forest plot    | 123 exact      | Yes      |
| F3     | Mechanism panel   | 162 exact      | Yes      |
| F4     | RQ1 held-out      | 94 exact       | Yes      |
| F5     | Occupancy-rooting | 67 exact       | Yes      |
| F6     | Case studies      | 425 exact      | Yes      |
| F7     | Self-test         | 0 (schematic)  | Yes      |

2. **Intensive margin is conditional, not causal.** The estimate conditions on  $Y \geq 1$  and is subject to selection bias.
3. **Field boundary not detected.** Wald test  $p = 0.638$ .
4. **MeSH retention slightly lower in the host-specificity test.** On MeSH, 12.4% of the  $A_{\text{cont}}$  effect is absorbed by field breadth.
5. **Audit reveals 292 MISSING\_IN\_REPORT rows.**
6. **Table 18 self-test: 4 of 11 claims NOT DONE.**
7. **Seeded-perturbation recall 0.60.**

## 6.9 What Was Learned

Iteration 5 refines the confirmed grafting finding along three dimensions. Both margins matter (IVW extensive +6.43 pp/SD [4.76, 8.11]; IVW intensive IRR/SD 1.183 [1.104, 1.267]). The effect is host-specific (IVW retention 0.969 [0.81, 1.10];  $A_{\text{lift}}$  IVW IRR/SD 1.34 [1.23, 1.46]). No field boundary is detected (Wald  $p = 0.638$ ). The emergence question’s primary closure is dead (Holm  $p = 0.147$ ); persistent-neighbour closure is confirmed (Holm  $p = 0.015$ ).

## 7 Run Bookkeeping

Table 39: Run trajectory. Ledger spend is cumulative LLM/tool/container spend and excludes compute-rental time.

| Iter. | Move   | Review score | Executed | Ledger spend (USD) |
|-------|--------|--------------|----------|--------------------|
| 1     | –      | –            | Yes      | –                  |
| 2     | –      | –            | Yes      | –                  |
| 3     | –      | –            | Yes      | –                  |
| 4     | –      | –            | Yes      | –                  |
| 5     | deepen | 5            | Yes      | \$0.105            |

The trajectory log records only the final iteration (Iteration 5) with explicit fields. All five iterations executed artifacts. The cumulative ledger spend through Iteration 5 was \$0.105 (LLM and tool costs only, excluding orchestrator compute-rental time).

Table 40: Run ledger: all artifacts across five iterations.

| Artifact                                     | Type       | Status           | Iter. |
|----------------------------------------------|------------|------------------|-------|
| gen_art_dataset_1 (pool)                     | dataset    | succeeded        | 1     |
| gen_art_dataset_2 (background)               | dataset    | failed (timeout) | 1     |
| gen_art_dataset_3 (MeSH)                     | dataset    | succeeded        | 1     |
| gen_art_dataset_4 (grounding)                | dataset    | succeeded        | 1     |
| gen_art_research_1 (dossier)                 | research   | succeeded        | 1     |
| gen_art_dataset_5 (hydrated)                 | dataset    | succeeded        | 2     |
| gen_art_experiment_1 (grounding)             | experiment | succeeded        | 2     |
| gen_art_experiment_2 (viability)             | experiment | succeeded        | 2     |
| gen_art_experiment_3 (emergence, main)       | experiment | succeeded        | 2     |
| gen_art_experiment_4 (emergence, MeSH)       | experiment | succeeded        | 2     |
| gen_art_evaluation_1 (audit)                 | evaluation | succeeded        | 3     |
| gen_art_experiment_5 (emergence deepen)      | experiment | succeeded        | 3     |
| gen_art_experiment_6 (breadth prediction)    | experiment | succeeded        | 3     |
| gen_art_experiment_7 (host-entry grafting)   | experiment | succeeded        | 3     |
| gen_art_experiment_8 (descriptive diffusion) | experiment | succeeded        | 3     |
| gen_art_evaluation_2 (grafting held-out)     | evaluation | succeeded        | 4     |
| gen_art_experiment_9 (MeSH replication)      | experiment | succeeded        | 4     |
| gen_art_evaluation_3 (closure held-out)      | evaluation | succeeded        | 4     |
| gen_art_evaluation_4 (descriptive held-out)  | evaluation | succeeded        | 4     |
| gen_art_evaluation_5 (adopter mechanism)     | evaluation | succeeded        | 4     |
| gen_art_evaluation_6 (K1/K3 tests)           | evaluation | succeeded        | 5     |
| gen_art_evaluation_7 (K2 test)               | evaluation | succeeded        | 5     |
| gen_art_evaluation_8 (full audit)            | evaluation | succeeded        | 5     |
| gen_art_research_2 (positioning dossier)     | research   | succeeded        | 5     |
| gen_art_evaluation_9 (paper figures)         | evaluation | succeeded        | 5     |

## 8 Hashes

## 9 What the Run Learned Overall and What Is Still Open

### 9.1 Summary of Confirmed Findings

**Diffusion: host-entry grafting is confirmed and replicated.** When a concept enters a new host subfield, the share of its initial co-occurrence partners that are native to that host predicts subsequent uptake by newcomer authors. The co-primary held-out fold gives IRR/SD 1.19 [1.06, 1.33] ( $p = 0.004$ ), MeSH replication gives IRR/SD 1.23 [1.12, 1.36] (Holm  $p < 0.001$ ), and the IVW pooled estimate is 1.26 [1.17, 1.36] with  $I^2 = 0$ . Co-transfer of origin companions is null across all folds and populations. The effect operates through both the extensive margin (IVW +6.43 pp/SD [4.76, 8.11]) and the intensive margin (IVW IRR/SD 1.183 [1.104, 1.267]), is host-specific rather than reflecting general accessibility (IVW retention 0.969), and does not vary by origin field (Wald  $p = 0.638$ ).

**Emergence: persistent-neighbour closure is confirmed.** The general closure deficit did not survive held-out confirmation (Holm  $p = 0.147$ ). Only persistent-neighbour closure, restricted to top-20 associates present in consecutive years, is confirmed ( $S = -1.073$ , Holm  $p = 0.015$ ).

Table 41: Pre-registration and specification hashes.

| Item                                 | SHA-256 prefix  |
|--------------------------------------|-----------------|
| Frozen concept frame                 | 80e3f244...0c44 |
| Viability pre-registration           | 5dd16d8a...5aac |
| Test population                      | 6a887fb4...5505 |
| R1a turnover pre-check spec          | 60f44c0f...     |
| Emergence deepen spec (SPEC3)        | 2f46d173...     |
| Host-entry grafting pre-registration | f800a0a9...     |
| Descriptive diffusion held-out spec  | 695166b4...     |
| Held-out confirmation spec (exp_5)   | 535c2dd3...     |
| Grafting held-out spec               | 8db17113...     |
| Discriminating-test spec             | a10b7f31...     |
| K1/K3 spec                           | 2435f909...     |
| K2 spec                              | 866c60a8...     |
| MeSH replication spec                | b7cdabf8...     |

Burt constraint replicates its positive sign (+0.105 [0.021, 0.187]), ruling out the simple brokerage explanation.

**Diffusion typology.** A  $k = 2$  typology separates localised ( $n = 136$ ) from broad-from-the-start ( $n = 66$ ) concepts, replicating on the held-out fold. Network expansion precedes disciplinary diffusion in 25 of 26 pooled main-arm dual-onset concepts.

**Adopter mechanism.** Adopters are 3.09 times more likely than matched non-adopters to have prior exposure to the concept’s entry partners (OR 3.09 [2.32, 4.28]), with enrichment strongest for origin-vocabulary (FOREIGN) partners (OR 2.88).

## 9.2 Dead Ends and Null Results

- Source-sink viability decomposition: blocked by insufficient citation-layer density (Gate A, 17.3%).
- Prediction of whether a concept will emerge: precursors add no value beyond baselines.
- Breadth prediction from openness: null on screen and held-out.
- Closure-anchoring link: null (pre-emergence closure does not predict later host-entry anchoring).
- Co-transfer: null in every specification and population.
- 3-channel diffusion typology: adds nothing beyond entropy trajectories.
- Community roles predicting host entry: null.
- Binary establishment threshold (EST\_bin): null on held-out.

## 9.3 Open Questions

- The population is dominated by physics, computer science, and biomedicine. Generalisation to social sciences and humanities is untested.

- The primary fully-saturated specification is underpowered. A larger concept pool would permit a stronger test.
- Whether the single-paper artefact concern is fully resolved: the pooled multi-team effect is positive, but the held-out alone is inconclusive.
- Whether the host-vocabulary gradient operates through genuine translation of ideas into host vocabulary or through artefacts of topic-classification overlap. The topic-score absorption test suggests less than 8% absorption, but a citation-independent venue control was not estimable (12% coverage).
- The absorptive-capacity vs topical-proximity ambiguity in the adopter mechanism.
- Whether persistent-neighbour closure is robust to alternative definitions of neighbourhood stability.